

assignment__02_assignment-02-b

September 17, 2020

Linear supervised regression

0. Import library

Import library

```
[34]: # Import libraries

# math library
import numpy as np

# visualization library
%matplotlib inline
from IPython.display import set_matplotlib_formats
set_matplotlib_formats('png2x', 'pdf')
import matplotlib.pyplot as plt

# machine learning library
from sklearn.linear_model import LinearRegression

# 3d visualization
from mpl_toolkits.mplot3d import axes3d

# computational time
import time
```

1. Load dataset

Load a set of data pairs $\{x_i, y_i\}_{i=1}^n$ where x represents label and y represents target.

```
[35]: # import data with numpy
data = np.loadtxt('profit_population.txt', delimiter=',')
```


2. Explore the dataset distribution

Plot the training data points.

```
[36]: x_train = data[:,0]
      y_train = data[:,1]
```

```
plt.title("Training data")
plt.xlabel("Population size (x 10$)")
plt.ylabel("Profit (x 10k$)")

plt.scatter(x_train, y_train , marker='x',color='r')
plt.show()
```



3. Define the linear prediction function

$$f_w(x) = w_0 + w_1x$$

0.0.1 Vectorized implementation:

$$f_w(x) = Xw$$

with

$$X = \begin{bmatrix} 1 & x_1 \\ 1 & x_2 \\ \vdots & \vdots \\ 1 & x_n \end{bmatrix} \quad \text{and} \quad w = \begin{bmatrix} w_0 \\ w_1 \end{bmatrix} \quad \Rightarrow \quad f_w(x) = Xw = \begin{bmatrix} w_0 + w_1x_1 \\ w_0 + w_1x_2 \\ \vdots \\ w_0 + w_1x_n \end{bmatrix}$$

Implement the vectorized version of the linear predictive function.

```
[44]: m = len(x_train)

# construct data matrix
X = np.append(np.ones((m, 1)), np.array([x_train]).T, axis=1)
# parameters vector
w = np.array([[1], [1]])
# predictive function definition

def f_pred(X,w):
    f = X.dot(w)
    return f

# Test predicitive function
y_pred = f_pred(X,w)
```

4. Define the linear regression loss

$$L(w) = \frac{1}{n} \sum_{i=1}^n \left(f_w(x_i) - y_i \right)^2$$

0.0.2 Vectorized implementation:

$$L(w) = \frac{1}{n} (Xw - y)^T (Xw - y)$$

with

$$Xw = \begin{bmatrix} w_0 + w_1x_1 \\ w_0 + w_1x_2 \\ \vdots \\ w_0 + w_1x_n \end{bmatrix} \quad \text{and} \quad y = \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_n \end{bmatrix}$$

Implement the vectorized version of the linear regression loss function.

```
[53]: # loss function definition
def loss_mse(y_pred,y):
    loss = ( (y_pred - y).T ).dot( y_pred - y ) / m
    return loss.item()

# Test loss function
y = np.array([y_train]).T # label
y_pred = f_pred(X, w) # prediction

loss = loss_mse(y_pred,y)
```

5. Define the gradient of the linear regression loss

0.0.3 Vectorized implementation: Given the loss

$$L(w) = \frac{1}{n}(Xw - y)^T(Xw - y)$$

The gradient is given by

$$\frac{\partial}{\partial w} L(w) = \frac{2}{n} X^T(Xw - y)$$

Implement the vectorized version of the gradient of the linear regression loss function.

```
[54]: # gradient function definition
def grad_loss(y_pred,y,X):
    global m
    grad = (2 / m) * (X.T).dot(y_pred-y)
    return grad

# Test grad function
y_pred = f_pred(X, w)
y = np.array([y_train]).T
grad = grad_loss(y_pred,y,X)
```

6. Implement the gradient descent algorithm

- Vectorized implementation:

$$w^{k+1} = w^k - \tau \frac{2}{n} X^T(Xw^k - y)$$

0.0.4 Implement the vectorized version of the gradient descent function.

0.0.5 Plot the loss values $L(w^k)$ with respect to iteration k the number of iterations.

```
[70]: # gradient descent function definition
def grad_desc(X, y, w_init, tau, max_iter):

    L_iters = [] # record the loss values
    w_iters = [] # record the parameter values
    w = w_init # initialization

    for i in range(max_iter): # loop over the iterations

        y_pred = f_pred(X, w) # linear prediction function
        grad_f = grad_loss(y_pred, y, X) # gradient of the loss
        w = w - (tau * grad_f) # update rule of gradient descent
        L_iters.append(loss_mse(y_pred,y)) # save the current loss value
        w_iters.append(w)
    return w, L_iters, w_iters
```

```

# run gradient descent algorithm
start = time.time()
w_init = np.array([[1], [1]])
tau = 0.003
max_iter = 20000

w, L_iters, w_iters = grad_desc(X,y,w_init,tau,max_iter)

print('Time=',time.time() - start) # plot the computational cost
print( L_iters[-1] ) # plot the last value of the loss
print( w_iters[-1] ) # plot the last value of the parameter w


# plot
plt.figure(2)
plt.plot(L_iters, color='b') # plot the loss curve
plt.xlabel('Iterations')
plt.ylabel('Loss')
plt.show()

```

```

Time= 0.1923840045928955
8.953942751950358
[[-3.89578088]
 [ 1.19303364]]

```

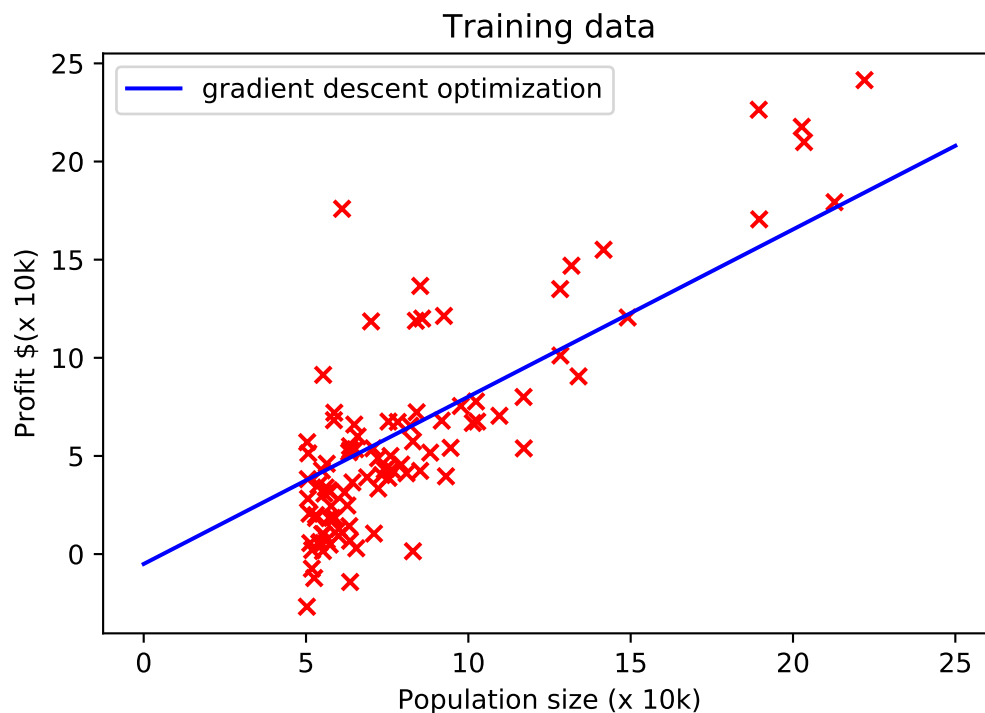


7. Plot the linear prediction function

$$f_w(x) = w_0 + w_1x$$

```
[64]: # linear regression model
x_pred = np.linspace(0,25,100) # define the domain of the prediction function
y_pred = w[0] + w[1] * x_pred # compute the prediction values within the given
    ↪ domain x_pred

# plot
plt.figure(3)
plt.scatter(x_train, y_train, color = 'r', marker='x')
plt.plot(x_pred, y_pred, color = 'blue', label = 'gradient descent
    ↪ optimization')
plt.legend(loc='best')
plt.title('Training data')
plt.xlabel('Population size (x 10k)')
plt.ylabel('Profit $(x 10k)')
plt.show()
```



8. Comparison with Scikit-learn linear regression algorithm

0.0.6 Compare with the Scikit-learn solution

```
[65]: # run linear regression with scikit-learn
start = time.time()
lin_reg_sklearn = LinearRegression()
lin_reg_sklearn.fit(np.array([x_train]).T, y) # learn the model parameters
print('Time=',time.time() - start)

# compute loss value
w_sklearn = np.zeros([2,1])
w_sklearn[0,0] = lin_reg_sklearn.intercept_
w_sklearn[1,0] = lin_reg_sklearn.coef_

print(w_sklearn)

loss_sklearn = loss_mse(f_pred(X, w_sklearn), y) # compute the loss from the
↳sklearn solution


print('loss sklearn=',loss_sklearn)
print('loss gradient descent=',L_iters[-1])

# plot
y_pred_sklearn = w_sklearn[0, 0] + w_sklearn[1, 0] * x_pred # prediction
↳obtained by the sklearn library

plt.figure(3)

plt.scatter(x_train, y_train, color = 'r', marker='x')
plt.plot(x_pred, y_pred, color = 'blue', label = 'gradient descent
↳optimization')
plt.plot(x_pred, y_pred_sklearn, color = 'orange', label = 'Scikit-learn
↳optimization')
plt.legend(loc='best')
plt.title('Training data')
plt.xlabel('Population size (x 10k)')
plt.ylabel('Profit $(x 10k)')
plt.show()
```

```
Time= 0.0008718967437744141
[[-3.89578088]
 [ 1.19303364]]
loss sklearn= 8.953942751950358
loss gradient descent= 11.05034451544502
```



9. Plot the loss surface, the contours of the loss and the gradient descent steps

```
[66]: # plot gradient descent
def plot_gradient_descent(X,y,w_init,tau,max_iter):

    def f_pred(X,w):

        f = X.dot(w)

        return f

    def loss_mse(y_pred,y):

        loss = ((y_pred - y).T).dot(y_pred - y) / m

        return loss

    # gradient descent function definition
    def grad_desc(X, y, w_init, tau, max_iter):

        L_iters = [] # record the loss values
        w_iters = [] # record the parameter values
        w = w_init # initialization
```



```

    for i in range(max_iter): # loop over the iterations

        y_pred = f_pred(X, w) # linear prediction function
        grad_f = grad_loss(y_pred, y, X) # gradient of the loss
        w = w - (tau * grad_f) # update rule of gradient descent
        L_iters.append(loss_mse(y_pred,y) ) # save the current loss value
        w_iters.append(w) # save the current w value

    return w, L_iters, w_iters

# run gradient descent
w, L_iters, w_iters = grad_desc(X, y, w_init, tau, max_iter)

# Create grid coordinates for plotting a range of  $L(w_0, w_1)$ -values
B0 = np.linspace(-10, 10, 50)
B1 = np.linspace(-1, 4, 50)

xx, yy = np.meshgrid(B0, B1, indexing='xy')
Z = np.zeros((B0.size,B1.size))

# Calculate loss values based on  $L(w_0, w_1)$ -values
for (i,j),v in np.ndenumerate(Z):
    Z[i,j] = loss_mse(f_pred(X, np.array([[xx[i][j]], [yy[i][j]]])), y)

# 3D visualization
fig = plt.figure(figsize=(15,6))
ax1 = fig.add_subplot(121)
ax2 = fig.add_subplot(122, projection='3d')

# Left plot
CS = ax1.contour(xx, yy, Z, np.logspace(-2, 3, 20), cmap=plt.cm.jet)
w_iters = np.array(w_iters)
ax1.scatter(w[0, 0], w[1, 0], color = 'r')
ax1.plot(w_iters[:, 0, 0], w_iters[:, 1, 0])

# Right plot
ax2.plot_surface(xx, yy, Z, rstride=1, cstride=1, alpha=0.6, cmap=plt.cm.
→jet)
ax2.set_zlabel('Loss  $L(w_0, w_1)$ ')
ax2.set_zlim(Z.min(), Z.max())

# plot gradient descent
Z2 = np.zeros([max_iter])

for i in range(max_iter):
    w0 = w_iters[i, 0, 0]
    w1 = w_iters[i, 1, 0]

```

```

Z2[i] = loss_mse( f_pred( X, np.array([[w0], [w1]])), y)

ax2.plot(w_iters[:, 0, 0], w_iters[:, 1, 0])
ax2.scatter( w[0,0], w[1,0], color = 'r')

# settings common to both plots
for ax in fig.axes:
    ax.set_xlabel(r'$w_0$', fontsize=17)
    ax.set_ylabel(r'$w_1$', fontsize=17)

```

```

[71]: # run plot_gradient_descent function
w_init = np.array([[1], [1]])
tau = 0.00001
max_iter = 100000

plot_gradient_descent(X,y,w_init,tau,max_iter)

```

